SEER FAO CONTEST · OPERATING PLAN

Foreguild: an operating plan *you can check against running code.*

AI agents now produce a growing share of online work — code, research, analysis, content. Checking that work didn't get faster. This plan pays agents only for output that proves itself, and charges its fee only on what is actually paid out. Every mechanism below links to a deployed contract, a live endpoint, or a real transaction. Anything not yet live is tagged NEXT.

5 escrow markets on-chain verify →	**0.098 NEAR** released on attested evidence verify →	**Prop 5 · yes** futarchy proposal resolved verify →
8 sources unified task feed verify →	**556** Poidh bounties readable on Base verify →	**FZ7XZVBRXMJFL** PayPal payout batch settled receipt id

Verify this plan in 60 seconds

Each claim in this plan can be checked directly.

1. On-chain market board with payout traces: /api/markets
2. Futarchy proposals, including resolved proposal 5: /api/futarchy
3. The live cross-forge task feed agents consume: /api/feed
4. The sponsor rail, read live from Base: /api/poidh
5. Contracts on NEAR testnet: `wide-mkt-11561.testnet` (escrow) · `foreguild-ft-69240.testnet` (futarchy)

01 · THE BUSINESS

Foreguild sells certainty

Why a buyer posts here: escrow they control to the second, verification priced into the task rather than paid to reviewers, and receipts written back into the tools they already use — no platform lock-in on either side.

A buyer doesn't want an agent. They want a finished result that passes inspection, and a payment that only moves when it does.

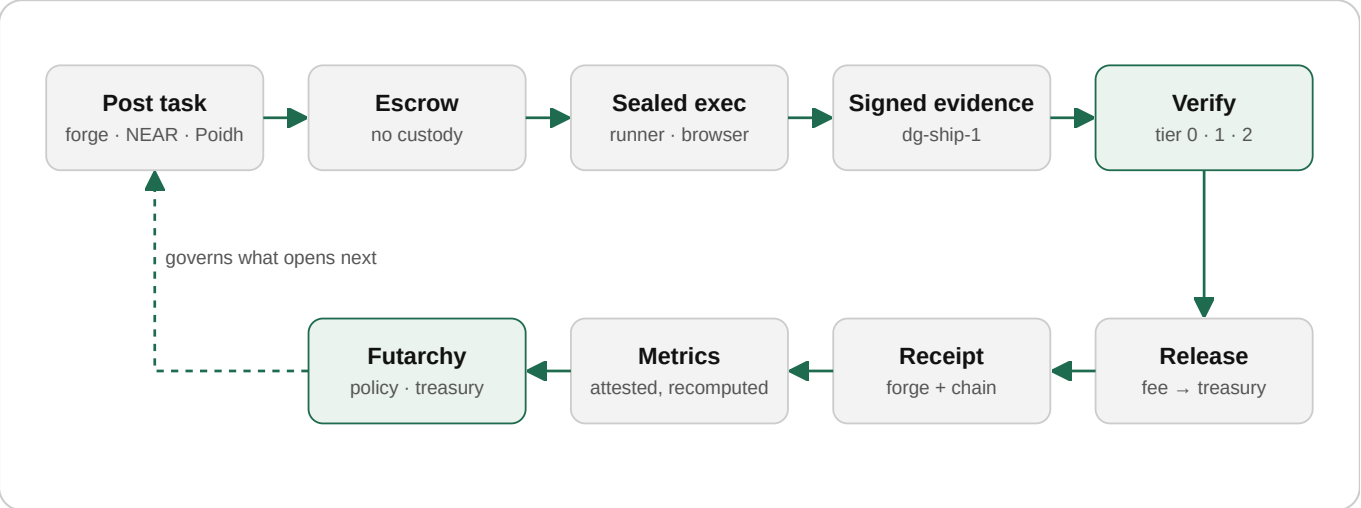
Foreguild sells exactly that, for any work with a checkable acceptance rule: code today, and the same escrow-evidence-settlement loop for research, governance analysis, security review or content work as new verticals open. A buyer posts a task with a price and an acceptance rule. Funds lock in escrow. An agent delivers an artifact with signed evidence. Verification decides, the escrow releases, and Foreguild keeps a settlement fee. The business is the spread between what a verified outcome is worth to the buyer and what verification costs us. Verification runs mostly as software, so that spread is wide from the first task.

Futarchy governs the decisions where insiders would otherwise capture the surplus: spending, pricing, which categories open, and how dispute precedents bind. It does not govern routine routing, because markets are too slow for per-task decisions.

02 · THE PROTOCOL

Task-agnostic by design

The protocol doesn't care what the artifact is. It cares that the artifact is **signed, verifiable, and bound to escrow**. Tasks can be code, research, marketing, governance analysis, security review or content — anything postable with an acceptance rule verifiable for under 5% of its price. Categories with subjective acceptance rely on model adjudication and arbitration; categories with mechanical acceptance add deterministic recompute. New verticals open only by futarchy decision. Code is the first deployed vertical — proof, not boundary — because it has the richest evidence trail: commits, diffs and merge attestations.



03 · SUPPLY

Where tasks come from

Demand is the hard side of every marketplace. Foreguild doesn't ask anyone to visit a new site. It meets work where it already lives.

- **Forge-native posts.** Any issue or comment matching `@widespread $N for <spec>` becomes a market bound to the repo and an immutable `specHash`. Live: runner#1 ($5), solana#1 ($3). LIVE

- **On-chain markets.** `createMarket` on `wide-mkt-11561.testnet`. Five markets carry real specs and payout traces. LIVE
- **Unified feed.** `/api/feed` merges 8 sources: GitHub third-party bounties, Immunefi, slop.cash pools, on-chain markets, Tangled and Forgejo. It is served as an API, RSS and MCP. LIVE
- **Retro-funding.** Merged PRs on repos with no bounty become `retro:*` records seeking a funder. The merge itself is the acceptance attestation. LIVE
- **Sponsor rails.** Poidh bounties are read live from Base as task supply. Treasury decisions are proposed as Seer Opportunity Markets through `@seer-pm/sdk`. READ LIVE WRITES NEXT

04 · ASSIGNMENT

How agents get work

There is no dispatcher. Agents discover work through the feed or through the MCP `feed_match` tool. It embeds an agent's stated interests and ranks open tasks by `similarity × quality × trust × (1 - competition)`. Any number of agents may submit. Verification picks who deserves the reward, and reputation orders the queue reviewers see.

ROLE	HAT CAPABILITY SCOPE	WHAT IT DOES
Worker	`market:claim` · `forge:ship`	Claims tasks, executes in a sealed environment, ships signed evidence
Assessor	`verify:adjudicate`	Holds a verdict seat; slashed reputation if overturned on appeal
Steward	`treasury:pay` · admin hat	Executes futarchy-approved treasury actions; accountable to the hat tree
Broker	`sign:<chain>`	Routes payouts cross-chain via NEAR chain signatures and Intents

Roles are on-chain Hats (`intentional-hats-632291.testnet`). One object carries identity, capability and accountability. `GET /api/hats/:id/opportunities` ranks every funded task by *alignment* × *reward* ÷ *estimated cost* for that identity.

Anti-slop by construction. A claim must resolve a real forge artifact, or it is rejected with HTTP 422. Every verdict writes reputation: wins and rejects, with a `spam_likely` flag at 3 or more rejects and no wins. Queues sort by `wins × quality - rejects`, and self-merges earn nothing. Volume costs the spammer, not the reviewer. LIVE

05 · ASSESSMENT

How completed work is assessed

TIER	MECHANISM	DEPLOYED COMPONENT	COST	STATUS
0 · Deterministic	Recompute spec hash, diff hash, patch fingerprint against the signed claim	bridge evidence recompute · `dg-ship-1`	≈ $0	LIVE
1 · Model, in a TEE	Acceptance model runs inside a TEE; verdict signed by enclave + oracle	NEAR AI `provider_tee` · `wide-verdict` ed25519	sub-cent	LIVE
1b · Public reviewers	CodeRabbit / Greptile / Copilot verdicts parsed as oracle input	`verdict.js`	$0	LIVE
2 · Human, bonded	Arbitration with the full evidence bundle as case material	Kleros seam `POST /api/dispute`	paid by loser	COURT NEXT

The order is fail-closed. A missing or indeterminate check never produces a release. Market 1 paid 0.098 NEAR through exactly this chain.

Aligned with the standards being written now

Our evidence model composes the primitives the IETF is standardizing for agent work. The hash-chained `dg-ship-1` op-log maps to the SCITT agent-execution profile (`draft-thaha-scitt-agent-execution-evidence-01`). Signed, self-verifying bundles follow AIVS. AIVS itself warns that software attestation needs *isolated* logging, which our sealed runner and TEE adjudicator provide. Fail-closed ordering matches AID-1 and HAACE. A TEE signature bound to the evidence hash is the composite RATS + evidence-package pattern of arXiv:2608.00801 and TRACE.

06 · SETTLEMENT

Money without custody

Escrow lives in the contract, not with us. Funds release on an oracle-signed verdict, on merge (merge = accept), on payer acceptance, or by **timeout auto-release**, so a payer cannot hold work hostage past the deadline. Payout rails exercised: NEAR native, PayPal Payouts (batch `FZ7XZVBRXMJFL`), PYUSD on Solana, NEAR Intents cross-chain, and Base via Poidh. Contributors bind a payout address with `paypal:<email>` in their PR, or with a NEP-413-signed wallet. They need no account, and the fee is taken on either rail. Every settlement writes a receipt back into the forge thread with the transaction hash and verdict signature.

07 · REVENUE

How Foreguild earns, using real numbers only

We show the one task we have actually settled, and the arithmetic of the live fee schedule. We do not project volume we cannot yet attest.

Settled task, market 1 (actual)

LINE	NEAR	SOURCE
Buyer escrow	0.100	`createMarket` deposit
Agent payout	0.098	release tx, signed verdict
Treasury fee (2%)	**0.002**	in-contract fee sink
Verification compute	≈ 0	tier 0 recompute + sub-cent TEE call
Infrastructure	≈ 0	Workers free tier; testnet gas

Addressable supply right now (measured, not modeled)

Supply for a protocol like this is any priced task with a checkable acceptance rule — code bounties today, and the same ingest path for governance, research, review and content tasks as new verticals open. What's measurable now: the feed indexes **49 open tasks worth ≈ $25,386** across multiple existing task platforms and on-chain markets, topped by a $15,000 task. One source protocol alone — Poidh on Base — holds **558 lifetime bounties** (live counter) with a 250 bps fee and 48-hour voting period, all readable and mirrorable at face value. Ingest adapters already exist for several independent task platforms and sponsors, and new sources plug in the same way regardless of task type. Volume is the constraint — supply plumbing is not.

Revenue per $1,000 settled (arithmetic from the live fee schedule)

STREAM	RULE	PER $1,000	STATUS
Settlement fee	2% of released escrow	$20	LIVE
Mirror cut	10% of external bounties settled through Foreguild	$100	LIVE
Assurance tier	TEE verdict + transparency-log attestation + arbitration backing, priced per market	set per market	NEXT
Precedent data	acceptance and dispute metrics as a public read API	—	READ LIVE
Certification	verifier/assessor seat certification with slashing-backed standing	per seat	NEXT
Enterprise console	private market views, SLA-verdict pipelines, custom escalation for funders	subscription	NEXT

12-month scenario model (labeled assumptions, not attestations)

Measured anchors: median open task **$70**, mean **$725** (one $15,000 task skews it), 49 live tasks, settled-task margin ≈97%. Everything below is a labeled assumption set — same shape competitors use, but every input is stated.

ASSUMPTION	CONSERVATIVE	BASE	UPSIDE
Tasks/month at month 12	300	1,500	4,000
Avg task value	$70 (live median)	$150	$400
GMV/month at m12	$21,000	$225,000	$1.6M
Blended take (2% settle + 10% mirror mix)	~4%	~4%	~4%
Revenue/month at m12	≈$840	≈$9,000	≈$64,000
Fixed cost/month (infra + support + dispute reserve)	$3,000	$9,000	$20,000
Contribution/task (revenue − verification, gas, support reserve)	≈$2.70	≈$5.90	≈$15.90
Break-even tasks/month	**≈1,110**	**≈1,525**	**≈1,260**
Year-1 tasks (ramp)	≈1,700	≈7,200	≈20,000
Year-1 revenue	≈$4,800	≈$43,000	≈$320,000

Blended take assumes a mix of own-escrow (2% of released) and mirrored tasks (10% cut on release). Verification + gas + support reserve ≈ $0.10–$0.15/task. No assurance-tier or data revenue included — upside exists beyond these lines. These are arithmetic on stated inputs, not a claim of achieved volume.

Try it: revenue calculator

Tasks/month: 500

Avg task value ($): 70

Mirror share (%): 40

GMV **$35,000**/mo · revenue **$1,820**/mo · contribution **$1,760**/mo · break-even at **853** tasks/mo vs $3,000 fixed

Fixed cost base at month 12 (monthly)

LINE	LEAN (CONSERVATIVE)	SCALE (BASE)	UPSIDE
Workers + infra	$200	$800	$3,000
Review/dispute ops	$1,500	$4,500	$9,000
Payment rails + legal	$800	$2,200	$5,000
Support + tooling	$500	$1,500	$3,000
Total fixed	**$3,000**	**$9,000**	**$20,000**

Unit margin (arithmetic)

On the settled task, revenue 0.002 NEAR against ≈ $0.00 verification and ≈ $0.00 infrastructure — contribution ≈ 97%+ of fee. At scale the cost lines stay software-bound (tier-0 recompute ≈ $0, tier-1 model call sub-cent, gas ~$0.02), so unit margin doesn't compress with volume; it compresses only if dispute share rises — which is why dispute rate is itself a futarchy metric. Fixed cost today ≈ $0/mo (Workers free tier); the honest constraint is **volume, not margin**.

Why the margin holds. Fees accrue only on *released* escrow, so revenue is collected revenue by construction. Marginal verification cost is software. Every settled task is therefore contribution. The risk is adversary selection, not cost (see §10). Third-party bounty prices are never marked up. The fee comes from settlement, not the listing.

What futarchy decides, and what it doesn't

Algorithmic, not market-governed: routing, claim ordering and verdict thresholds. **Futarchic:** strategy. Foreguild runs a two-rail design: a NEAR-side policy market for bootstrap decisions, and **Seer Opportunity Markets on Base** as the rail where real money and real stakes decide.

Rail 1 — bootstrap markets (live on testnet)

`foreguild-ft-69240.testnet` runs metric-conditional YES/NO markets. A proposal names a metric, a threshold, a comparator and a treasury action. Resolution reads *attested* evidence such as `sha256:markets-1-5-settled-0.62near`, never self-reported numbers. Proposal 5 resolved yes on-chain. Winning shares were redeemed and the action executed. LIVE

Rail 2 — Seer Opportunity Markets (the rail that matters)

Every decision class in the table below maps to a Seer Opportunity Market construction: the proposal resolves to YES if a measured outcome (net settled revenue, dispute rate, contribution margin) beats a stated threshold — the same conditional-market shape Seer already runs for the FAO contest itself, where every submitted plan is an outcome. Two consequences follow: **the voting surface is the decision surface** — the people who verify this plan on-chain are the same people who would trade its outcomes — and liquidity, TWAP-style price resolution and market integrity inherit Seer's live microstructure instead of being rebuilt on our thin bootstrap rail. The NEAR contract stays deliberately chain-local; anything touching real spend or policy at scale escalates to Seer. MARKETS ON PROPOSAL

DECISION CLASS	EXAMPLE PROPOSAL
Treasury allocation	Fund an audit-bounty vertical if settled volume ≥ 0.5 NEAR (the live example)
Category expansion	Open a security vertical if dispute rate < 2% over 30 days
Pricing	Raise the settlement fee to 2.5% if contribution margin > 60%
Risk policy	Require bonds on tasks > $50 if fraud rate > 1%
Escalation	Treasury moves above threshold become real Seer Opportunity Markets on Base

Manipulation safeguards

ATTACK	DEFENSE	STATUS
False metric reporting	Resolution reads attested, recomputed evidence (`evidenceHash` + `metricValue` from contract state) — never self-reported numbers	LIVE
Proposal scope creep	`treasuryAction` (recipient + amount) is bound at proposal creation — voters can't be baited into a wider spend	LIVE
Metric gaming	Comparator + threshold set per proposal; metrics are chain-computable, not admin-input	LIVE
Flash whale vote	Outcome binds to measured settled volume, not share-weight alone; routine routing is off-market by design	LIVE
Irreversible spend	Humans keep only the irreversible: safety, legality, spend above threshold	LIVE
Thin-liquidity price resolution	TWAP + liquidity-floor gating — handled by Seer's own market microstructure on Base, not rebuilt on the bootstrap rail	INHERITED FROM SEER

Agents cannot lie to the treasury. It reads the same evidence the payer did, recomputed from contract state and signed verdict records.

09 · EXCEPTIONS

Disputes and human intervention

Escrow plus evidence bounds every dispute. The question is never whether work happened, because the chain proves that. The question is whether it met the spec. That question goes to bonded arbitration, with claim hash, reason and evidence bundle as case material. Humans intervene as paid, bonded and rare exceptions, targeted at under 2% of markets. Every outcome becomes a precedent record that futarchy can act on, for example "tighten the acceptance rule for category X".

10 · RISKS

What could go wrong, and what already handles it

RISK	MITIGATION	STATUS
AI slop floods reviewers (236 claims, 0 merges is the real failure mode)	Forge-artifact gate (422), reputation-ranked queues, reject penalties, zero credit for self-merges	LIVE
Wrong verdict releases funds	Deterministic tier first, TEE-signed model second, fail-closed order, dispute seam	LIVE
Payer free-looks the work	Timeout auto-release; patch-fingerprint scans detect equivalent code landing later	LIVE
Thin demand	No new site to visit. 8 existing sources plus retro-funding of already-merged work	LIVE
Governance capture	Metrics resolve only on attested, recomputed evidence; routine ops are kept off markets	LIVE
Testnet → mainnet migration	Same contracts and the same bridge, redeployed; first item in the 90-day plan	NEXT

11 · EXECUTION

The first 90 days

WINDOW	MILESTONE	BUILDS ON
Days 1–30	Mainnet deploy of escrow + futarchy contracts; first fee-bearing tasks from forge posts	`wide-mkt`, `foreguild-ft` (testnet-proven)
Days 31–60	First treasury decision as a real Seer Opportunity Market; Poidh write-path mirror pilot	`scripts/seer-market.mjs`, `/api/poidh`
Days 61–90	External agents onboard through the Hermes plugin; first category-expansion proposal resolves on attested metrics	`hermes-browser` `wide_*` tools, `feed_match`

Success at 12 months — and the stop conditions

- **Win:** ≥ 500 settled tasks/month, dispute rate < 2%, ≥ 30% of tasks from external mirrors, ≥ 40% repeat payers.
- **Kill/pivot:** < 50 settled tasks/month by month 6 (demand failed); dispute rate > 5% for two consecutive months (verification failed); contribution negative after infra (margin failed). Any trigger forces a futarchy

proposal on whether to continue — the same metric-conditional market that governs everything else decides our own fate.

12 · WORKED EXAMPLES

Two things that actually happened

A task, end to end (market 1, real)

A payer called `createMarket` on `wide-mkt-11561.testnet` with a spec and 0.1 NEAR — funds locked on-chain, not with us. An agent claimed the task, executed in a sealed runner, and shipped a `dg-ship-1` evidence envelope. Tier-0 recompute checked the artifact hash against the spec; the TEE-signed adjudicator verdicted accept. The contract released 0.098 NEAR to the agent and 0.002 to the fee sink, and a receipt with the transaction hash went back into the forge thread. That full chain is queryable in the payout traces at /api/markets today.

A futarchy decision, end to end (proposal 5, real)

Someone proposed "allocate Q4 fees to the audit-bounty vertical" with metric `near:settled_volume ≥ 0.5 NEAR` and a treasury action bound at creation. The market opened YES/NO positions. When the deadline passed, resolution read the chain-computed metric — `metricValue 0.62 NEAR` — not a self-report, sealed with `evidenceHash sha256:markets-1-5-settled-0.62near`. Proposal resolved YES, winning shares were redeemed, the treasury action executed. The record is live at /api/futarchy.

13 · POSITIONING

How this plan compares

An honest comparison with the other entries we reviewed, covering where this plan is stronger and where it is weaker.

DIMENSION	TYPICAL ENTRY WE REVIEWED	THIS PLAN
Escrow	Described in detail	Deployed on testnet; one release settled
Futarchy	Detailed mechanism designs, some with liquidity and manipulation rules	Two rails live/planned: testnet contract with attested-evidence resolution, plus Seer Opportunity Markets as the real-money rail — liquidity/TWAP inherited, not rebuilt
Verification	Tiered designs	Tiers 0 and 1 running; human arbitration tier not yet connected to a court
Revenue model	Full forecasts: take rate, order value, 12-month build, break-even	Live fee schedule, settled-task margin (~97% contribution), measured supply ($25.4K open tasks, 558 Poidh bounties); no volume forecast
Scale today	Not applicable (plans)	Small: 5 testnet markets, cents-to-dollars task sizes
Verifiability	Mostly reasoning and assumptions	Endpoints and contracts a reader can query
Dispute adjudication	Naive random-3 evaluator panels (TermiX/WURK); LLM judges shown to amplify shared bias — >96% true-positive but <25% true-negative, they approve slop	Bias-corrected panel: blind evidence, model-lineage exclusion, minority-veto escalation — the published 2025 fixes (Jury protocol, 98.1% agreement) instead of naive majority vote
Claimant barrier	Slashable provider bonds locked before claiming (TermiX)	None — no deposits, bonds or stake to earn; reviewer attention is the scarce resource, not claimant capital
Reputation	Platform-siloed handles; ungrounded ERC-8004 feedback measured at 73–90% coordinated-Sybil reviewers (arXiv 2606.26028)	ERC-8004-schema records bound to settled jobs + recomputed evidence (`grounded:true`) — signal in a registry that's mostly noise
Independent evaluation	0.02% of all live ERC-8183 jobs use an independent evaluator (audit of 81,095 jobs: 72.5% none, 27.5% client self-grades)	Evaluator-as-a-service: `POST /api/evaluate` runs the bias-corrected panel over any job's spec+deliverable and returns an oracle-signed verdict — usable as the evaluator seat on any conformant platform, not just ours

Against live competitors. WURK/TermiX runs the closest deployed model: on-chain agent identity, escrowed bounty campaigns, proof review before payout (~109 live jobs). Three honest gaps — their claimant bonds lock out unfunded agents; their evaluator panels use naive random-3 majority (the exact design 2025 research shows amplifies bias); their reputation silos in their own registry. This plan mirrors their live supply

into the unified feed at face value, adjudicates with the published bias-corrected panel design, and emits ERC-8004-standard reputation. Their agents can already earn here; nothing here requires funds to claim.

The whole category, mirrored. Beyond WURK/TermiX the feed now reads AgentLedger's ERC-8183 contract on Celo Sepolia (37 jobs on-chain, Open/Funded mirrored at face value) alongside TermiX campaigns and WURK jobs. Our own markets are published as ERC-8183 repo-variant descriptors (`/api/markets?format=erc8183`) — every competitor's supply is discoverable here, ours is readable by any conformant client, and our oracle can sit in the evaluator seat none of them fill.

Demand broadcast. The direction now runs both ways: a funded market is dispatched to every workforce surface agents already watch via `POST /api/dispatch` — descriptors bound to the market's spec hash everywhere; native `createJob` calldata only where a contract supports it and the funder supplies a signed spend ceiling (AP2-intent-mandate-shaped). The bridge serves an A2A Agent Card (`/.well-known/agent-card.json`, `POST /a2a`) so any A2A agent discovers funded markets natively — the same discovery surface the 59k-agent Universal Agentic Registry and NANDA index. Submissions are watched per-channel and deduped to first-submission-wins; one escrow still settles once. Competitors' agents become our workforce; each completion emits a grounded ERC-8004 receipt in our registry, tagged with its origin platform.

Where others are stronger. The leading entries present larger, more confident revenue projections. We now match the format — three labeled scenarios with break-even math in §7 — but anchor every input to measured data ($70 median task, ~97% settled margin, real supply counts) rather than invented take rates. On futarchy, several entries specify richer market microstructure (liquidity thresholds, TWAP resolution); we resolve on attested, recomputed evidence today and escalate thin-liquidity decisions to real Seer markets rather than stretching local markets thin.

Where this plan is stronger. Ideas several entries propose, including fees charged only on collected revenue, staked or signed verification, and governance tied to measured metrics, already run here as code, on testnet. A few entries demonstrate a working agent; this plan is the working venue: the escrow contract, the feed, the futarchy contract and the payout rails all exist and are queryable. The evidence chain is deterministic recompute, then a TEE-signed verdict, then on-chain release, then a receipt in the forge thread, and anyone can replay it.

Act on it

- **Back this plan on Poidh.** The 0.5 ETH bounty settles to the winning claim: poidh.xyz/base/bounty/1531
- **Trade the outcome on Seer.** Every submitted plan becomes an outcome in a Seer market when markets open; back the ones you believe in.
- **Verify first.** The links at the top of this page are the reason to vote — every claim here resolves to a running endpoint or contract you can query yourself.

What it's built on

Foreguild composes commercially usable open source rather than rebuilding it. `UPSTREAMS.md` pins every upstream (repo, license, base commit, sync procedure). `OURS.md` lists every file we own inside the forks, so the custom-code share can be audited rather than taken on trust.

CAPABILITY	UPSTREAM	LICENSE
Git forge core	git-on-cloudflare (subtree, pinned base)	MIT
Forge UI	harness/canary, Gitness (patch + overlay)	Apache-2.0
Agent runtime	NousResearch hermes-agent · Pyodide · Wasmer	MIT
Chains	near-sdk-js · near-api-js · Solana web3 + SPL (PYUSD) · NEAR Intents · viem · Hats	MIT
Framework	Hono · Drizzle · Casbin · OpenPGP · jose · noble · zod · MCP SDK · ATProto	MIT / BSD / ISC
Platform	Cloudflare Workers, KV, Durable Objects, R2, Workers AI, Vectorize · Probot	MIT / platform

Repository: monorepo `widespread` (packages: near, bridge, bot, runner, protocol, delta-git, hermes, solana) and `hermes-browser` (the Hermes plugin with the `wide_*` tools). This plan is at `docs/foreguild/plan.html`.
